

master-dataset.ipynb

Scope: Construct the working dataset based on three files: FanGraphs-Pitching-Standard.csv, FanGraphs-Pitching-Advanced.csv, and FanGraphs-Pitching-Value.csv. Export the created dataset to a .csv file for us in WAR-Predictor.ipynb.

```
In [1]: import pandas as pd
import os
```

```
In [3]: #Reading FanGraphs data into Pandas dataframes
pitching_standard = pd.read_csv('FanGraphs-Pitching-Standard.csv')
pitching_advanced = pd.read_csv('FanGraphs-Pitching-Advanced.csv')
pitching_value = pd.read_csv('FanGraphs-Pitching-Value.csv')
```

```
In [4]: #Removing irrelevant statistics from pitching_value
pitching_value = pitching_value.drop(columns=["RA9-WAR", "BIP-Wins", "LOB-Wins", "FDP-Wins", "RAR", "Dollars"])

#Joining datasets to create master dataset
pitching_mds = pd.merge(pitching_standard, pitching_advanced, on=["playerid", "Season", "Name", "Team"])
pitching_mds = pd.merge(pitching_mds, pitching_value, on=["playerid", "Season", "Name", "Team"])

#Removing duplicate columns
pitching_mds = pitching_mds.drop(columns=["ERA_y"])
pitching_mds = pitching_mds.rename(columns={"ERA_x" : "ERA"})

#Export master dataset to csv file
pitching_mds.to_csv('pitching-masterdataset.csv', index=False)
```

```
In [ ]:
```

WAR=Predictor.ipynb

Scope: Create a multiple linear regression model and random forest regression model to predict a pitcher's WAR based on other statistics.

```
In [1]: #Package Imports
import pandas as pd
import numpy as np
import sklearn.metrics as metrics
import os
import matplotlib.pyplot as plt
from sklearn.linear_model import LinearRegression
from sklearn.model_selection import train_test_split
from sklearn.metrics import r2_score
from sklearn.ensemble import RandomForestRegressor

In [2]: #Load master dataset
df_pitching = pd.read_csv('pitching-masterdataset.csv')

print(df_pitching.describe())

count    Season      W      L      ERA      G  \
mean    2015.751023  4.230567  5.939801  3.900009  40.379018
std      3.072060  4.261957  3.654775  1.175255  20.029821
min     2011.000000  0.000000  0.000000  0.540000  8.000000
25%     2013.000000  3.000000  3.000000  3.100000  25.000000
50%     2016.000000  5.000000  5.000000  3.810000  33.000000
75%     2018.000000  9.000000  8.000000  4.607500  61.000000
max     2021.000000  24.000000  19.000000  9.580000  85.000000

count    GS      CG      SHO      SV      WLD ...  \
mean    13.345120  0.274693  0.134424  3.141146  4.746055 ...
std     12.577148  0.759965  0.443922  8.742923  7.737213 ...
min      0.000000  0.000000  0.000000  0.000000  0.000000 ...
25%      0.000000  0.000000  0.000000  0.000000  0.000000 ...
50%     12.000000  0.000000  0.000000  0.000000  0.000000 ...
75%     26.000000  0.000000  0.000000  1.000000  8.000000 ...
max     35.000000  11.000000  6.000000  57.000000  41.000000 ...

count    WHIP      BABIP      ERA-      FIP-      xFIP-  \
mean    1.277022  0.289800  94.783168  96.512858  97.558153
std      0.200009  0.032317  27.547303  21.260601  16.754184
min      0.550000  0.170000  13.000000  21.000000  23.000000
25%      1.140000  0.269000  76.000000  83.000000  87.000000
50%      1.270000  0.291000  94.000000  97.000000  98.000000
75%      1.410000  0.311000  111.000000  110.000000  109.000000
max      2.090000  0.388000  220.000000  178.000000  158.000000

count    FIP      E-F      xFIP      SIERA      WAR
mean    3.967896  -0.067987  4.009936  3.896634  1.287084
std      0.911166  0.747558  0.722406  0.771167  1.404701
min      0.780000  -2.400000  0.880000  0.760000  -1.600000
25%      3.360000  -0.550000  3.550000  3.400000  0.300000
50%      3.950000  -0.110000  4.030000  3.930000  1.000000
75%      4.540000  0.400000  4.480000  4.430000  1.900000
max      7.650000  2.880000  6.530000  6.280000  9.000000

[8 rows x 39 columns]

In [3]: #Removing & signs from fields that contain them
for x in range(len(df_pitching)):
    df_pitching['K%'][x] = df_pitching['K%'][x][:~1]
    df_pitching['BB%'][x] = df_pitching['BB%'][x][:~1]
    df_pitching['K-BB%'][x] = df_pitching['K-BB%'][x][:~1]
    df_pitching['LOB%'][x] = df_pitching['LOB%'][x][:~1]

#Split the master dataset into training set and testing set
df_training, df_testing = train_test_split(df_pitching, test_size=0.25, random_state=42, shuffle=True)

print("Number of records in training set: " + str(len(df_training)))
print("Number of records in testing set: " + str(len(df_testing)))
print("Total number of records in master dataset: " + str(len(df_pitching)))

/var/folders/g0/4lrv_ph94r7ljngnydc3cl840000gn/t/ipykernel_13369/3217800170.py:3: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
    df_pitching['K%'][x] = df_pitching['K%'][x][:~1]
/var/folders/g0/4lrv_ph94r7ljngnydc3cl840000gn/t/ipykernel_13369/3217800170.py:4: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
    df_pitching['BB%'][x] = df_pitching['BB%'][x][:~1]
/var/folders/g0/4lrv_ph94r7ljngnydc3cl840000gn/t/ipykernel_13369/3217800170.py:5: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
    df_pitching['K-BB%'][x] = df_pitching['K-BB%'][x][:~1]
/var/folders/g0/4lrv_ph94r7ljngnydc3cl840000gn/t/ipykernel_13369/3217800170.py:6: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
    df_pitching['LOB%'][x] = df_pitching['LOB%'][x][:~1]
Number of records in training set: 2566
Number of records in testing set: 856
Total number of records in master dataset: 3422

In [4]: #Determine x and y variables. Let x represent the independent variables and y represent the dependent variable.
y_train = df_training['WAR']
X_train = df_training.drop(columns=['WAR', 'Season', 'Name', 'Team', 'playerid'])

y_test = df_testing['WAR']
X_test = df_testing.drop(columns=['WAR', 'Season', 'Name', 'Team', 'playerid'])

In [5]: #Construct multiple regression model
regr = LinearRegression()
regr.fit(X_train, y_train)

#Apply multiple linear regression model prediction
y_pred_linear = regr.predict(X_test)

#R-squared results to determine the success of the multiple linear regression model
r2_linear = r2_score(y_test, y_pred_linear)
print("Multiple Linear Regression R-Squared Score: " + str(r2_linear))

Multiple Linear Regression R-Squared Score: 0.9747554731866983

In [6]: #Construct random forest regression model
forest_regr = RandomForestRegressor(n_estimators=100, random_state=0)
forest_regr.fit(X_train, y_train)

#Apply random forest model prediction
y_pred_forest = forest_regr.predict(X_test)

In [7]: def regression_metrics(y_test, y_pred):
    #Regression metrics
    explained_variance = metrics.explained_variance_score(y_test, y_pred)
    mean_absolute_error = metrics.mean_absolute_error(y_test, y_pred)
    mean_squared_error = metrics.mean_squared_error(y_test, y_pred)
    median_absolute_error = metrics.median_absolute_error(y_test, y_pred)

    #Output regression metrics
    print("Explained Variance: " + str(explained_variance))
    print("Mean Absolute Error: " + str(mean_absolute_error))
    print("Mean Squared Error: " + str(mean_squared_error))
    print("Median Absolute Error: " + str(median_absolute_error))

In [9]: #Regression metrics for Linear Regression model
print("Multiple Linear Regression Statistics:")
regression_metrics(y_test, y_pred_linear)

print("-----")

#Regression metrics for Random Forest Model
print("Random Forest Model Statistics:")
regression_metrics(y_test, y_pred_forest)

Multiple Linear Regression Statistics:
Explained Variance: 0.9747554731866983
Mean Absolute Error: 0.16817630745819503
Mean Squared Error: 0.048794760547809166
Median Absolute Error: 0.13078011257812044
-----
Random Forest Model Statistics:
Explained Variance: 0.9701821411621511
Mean Absolute Error: 0.17526985981308413
Mean Squared Error: 0.05766230490654204
Median Absolute Error: 0.131
```

In []: